

Functional programming in Scala

A practical introduction

Guillaume Bogard - guillaumebogard.dev

Goals

At the end of this training session, you will:

- Be able to write complete Scala programs on your own
- Have a solid understanding of some of the most important concepts surrounding functional programming
- Be able to handle asynchronous communication on TCP sockets

What we will build

- A command-line Tic Tac Toe game that can be played both locally and on the local network
- We'll follow a *purely functional* style (no mutable state, strict encapsulation of side-effects)
- The game logic will be properly tested
- We'll use common Scala libraries

What is Scala?

- Scala was born in 2003 at the EPFL
- It's a programming language that mixes ideas from strongly typed functional programming languages such as Haskell and object-oriented languages such as Java
- It runs mostly on the JVM but also in the browser through Scala.js and natively through Scala Native
- On the JVM, it is completely interoperable with Java and its massive ecosystem

Scala in the wild

- Scala is a general purpose language. Its main use cases are
 - Web development
 - Data engineering
 - Highly concurrent applications more generally
- It is the most widely used functional programming language today
- It is used in production by Twitter, Zalando, Netflix, Disney, Fortnite, LinkedIn ...

Installing your development environment

To productive in Scala you'll need two things :

- `sbt`, which you can install through Homebrew or Sdkman
 - `sdk install sbt`
- A Scala IDE : Either IntelliJ IDEA or a text editor with the *Metals* extension
 - To use *Metals* you'll also need a JDK 8 installed

Meet sbt

Sbt is the *de facto* build tool for Scala.

- It defines the metadata of your Scala project
- It fetches your external dependencies
- It leverages the Scala compiler to build your sources into .class files
- It runs your tests
- It allows you to define custom tasks and run them

It's time to open your first Scala project
Checkout the exercises branch and open the project

Hello, World!

- In Scala, the simplest Hello World is written as follows:

```
object Main extends App {  
  println("Hello world")  
}
```

- To run the project, we do

```
sbt run
```

Functions and expressions

Functions.scala

Expressions

- An expression is a combination of terms that can be reduced to a value, e.g.:
 - `45 + 3`
 - `"Foo Bar Baz"`
 - `"Hello" + " " + "World"`
- Named values are defined using the `val` keyword
 - `val a = 12`
- Named values are immutable
- Every value has a type

Typing

- All values have a type, including functions
- Types in Scala are organized into a hierarchy
 - Children of `AnyVal` : primitive value types such as `Boolean`, `Unit`, `String`, `Int`...
 - Children of `AnyRef` : object types such as `List`, `Option` and user-defined classes
 - `Any` is the root of all types

Type inference

- The type of a value can in many cases be inferred ("guessed") by the compiler, meaning:
 - You can define it explicitly:
`val name: String = "George"`
 - Or let the compiler do its work:
`val name = "George"`

Complex expressions

- Scala is an expressive language : it doesn't make a distinction between expressions and statements that imperative languages make
- In other words : every construct in Scala yields a value, including control structure such as if/else
- Curly braces are use to delimitate complex expressions
- The effective value of a complex expression is the last line of the block

Functions

- Methods in Scala are defined with the `def` keyword like so:
`def square(value: Int) = value * value`
- The return type of a method can be defined explicitly:
`def square(value: Int): Int = value * value`

Methods and functions are slightly different things in Scala, but for now we'll use them interchangeably.

Different evaluation strategies

Evaluation is the process by which *expressions* are reduced to *values*

Scala supports different evaluation strategies for named values, methods, and methods arguments.

`val` **vs** `def`

- `vals` are evaluated the moment they are defined
- `defs` are evaluated every time they are accessed
- `lazy vals` are evaluated when they are first accessed and then memoized for later use

val vs def, an example

- This would crash immediately:

```
val a: String = throw new Exception("BOOM")  
println("Hello")  
println(a)
```

- This would print "Hello" then crash:


```
def a: String = throw new Exception("BOOM")  
println("Hello")  
println(a)
```


Call-by-name, call-by-value

- By default, method arguments are evaluated before the body of the function. We call it *call-by-value*
- *Call-by-name* arguments are evaluated only when accessed within the body of the function, not before
 - It means we can partially evaluate the body of a function without evaluating its arguments
 - They are defined with a fat arrow like so

```
def greet(name: => String) = println("Hello " + name)
```

**What is the billion
dollar mistake? 💰**



I call it my billion-dollar mistake. It was the invention of the null reference in 1965. [...] This has led to innumerable errors, vulnerabilities, and system crashes, which have probably caused a billion dollars of pain and damage in the last forty years.

– Sir Tony Hoare

Options

- Options are designed to avoid the infamous `NullPointerException`
- Options can be one of two things : Some or None
- They make the code more explicit, remove the need to search in documentation for possible nulls and prevents runtime exceptions in many cases
- They are explicitly checked by the compiler

Building options

- To build an empty option, use the None constructor:

```
val batmanFather: Option[String] = None
```

- To build an option from a value, use the Some constructor:

```
val batgirlFather: Option[String] = Some("Jim Gordon")
```


Working with options

- `map`
 - If you have `Option[A]` and `A => B`, then you get `Option[B]`
- `flatMap`
 - If you have `Option[A]` and `A => Option[B]`, then you get `Option[B]`
- `flatten`
 - If you have `Option[Option[A]]`, then you get `Option[A]`
- `getOrElse`
- `isEmpty`

Case classes

Case classes are used to create immutable values out of multiple named fields. Think of them as immutable structs.

```
case class Position(x: Int, y: Int)
```

Once the case class is defined, instances can be created like so :

```
val playerPosition = Position(18, 4)
```

- Case classes fields are public:

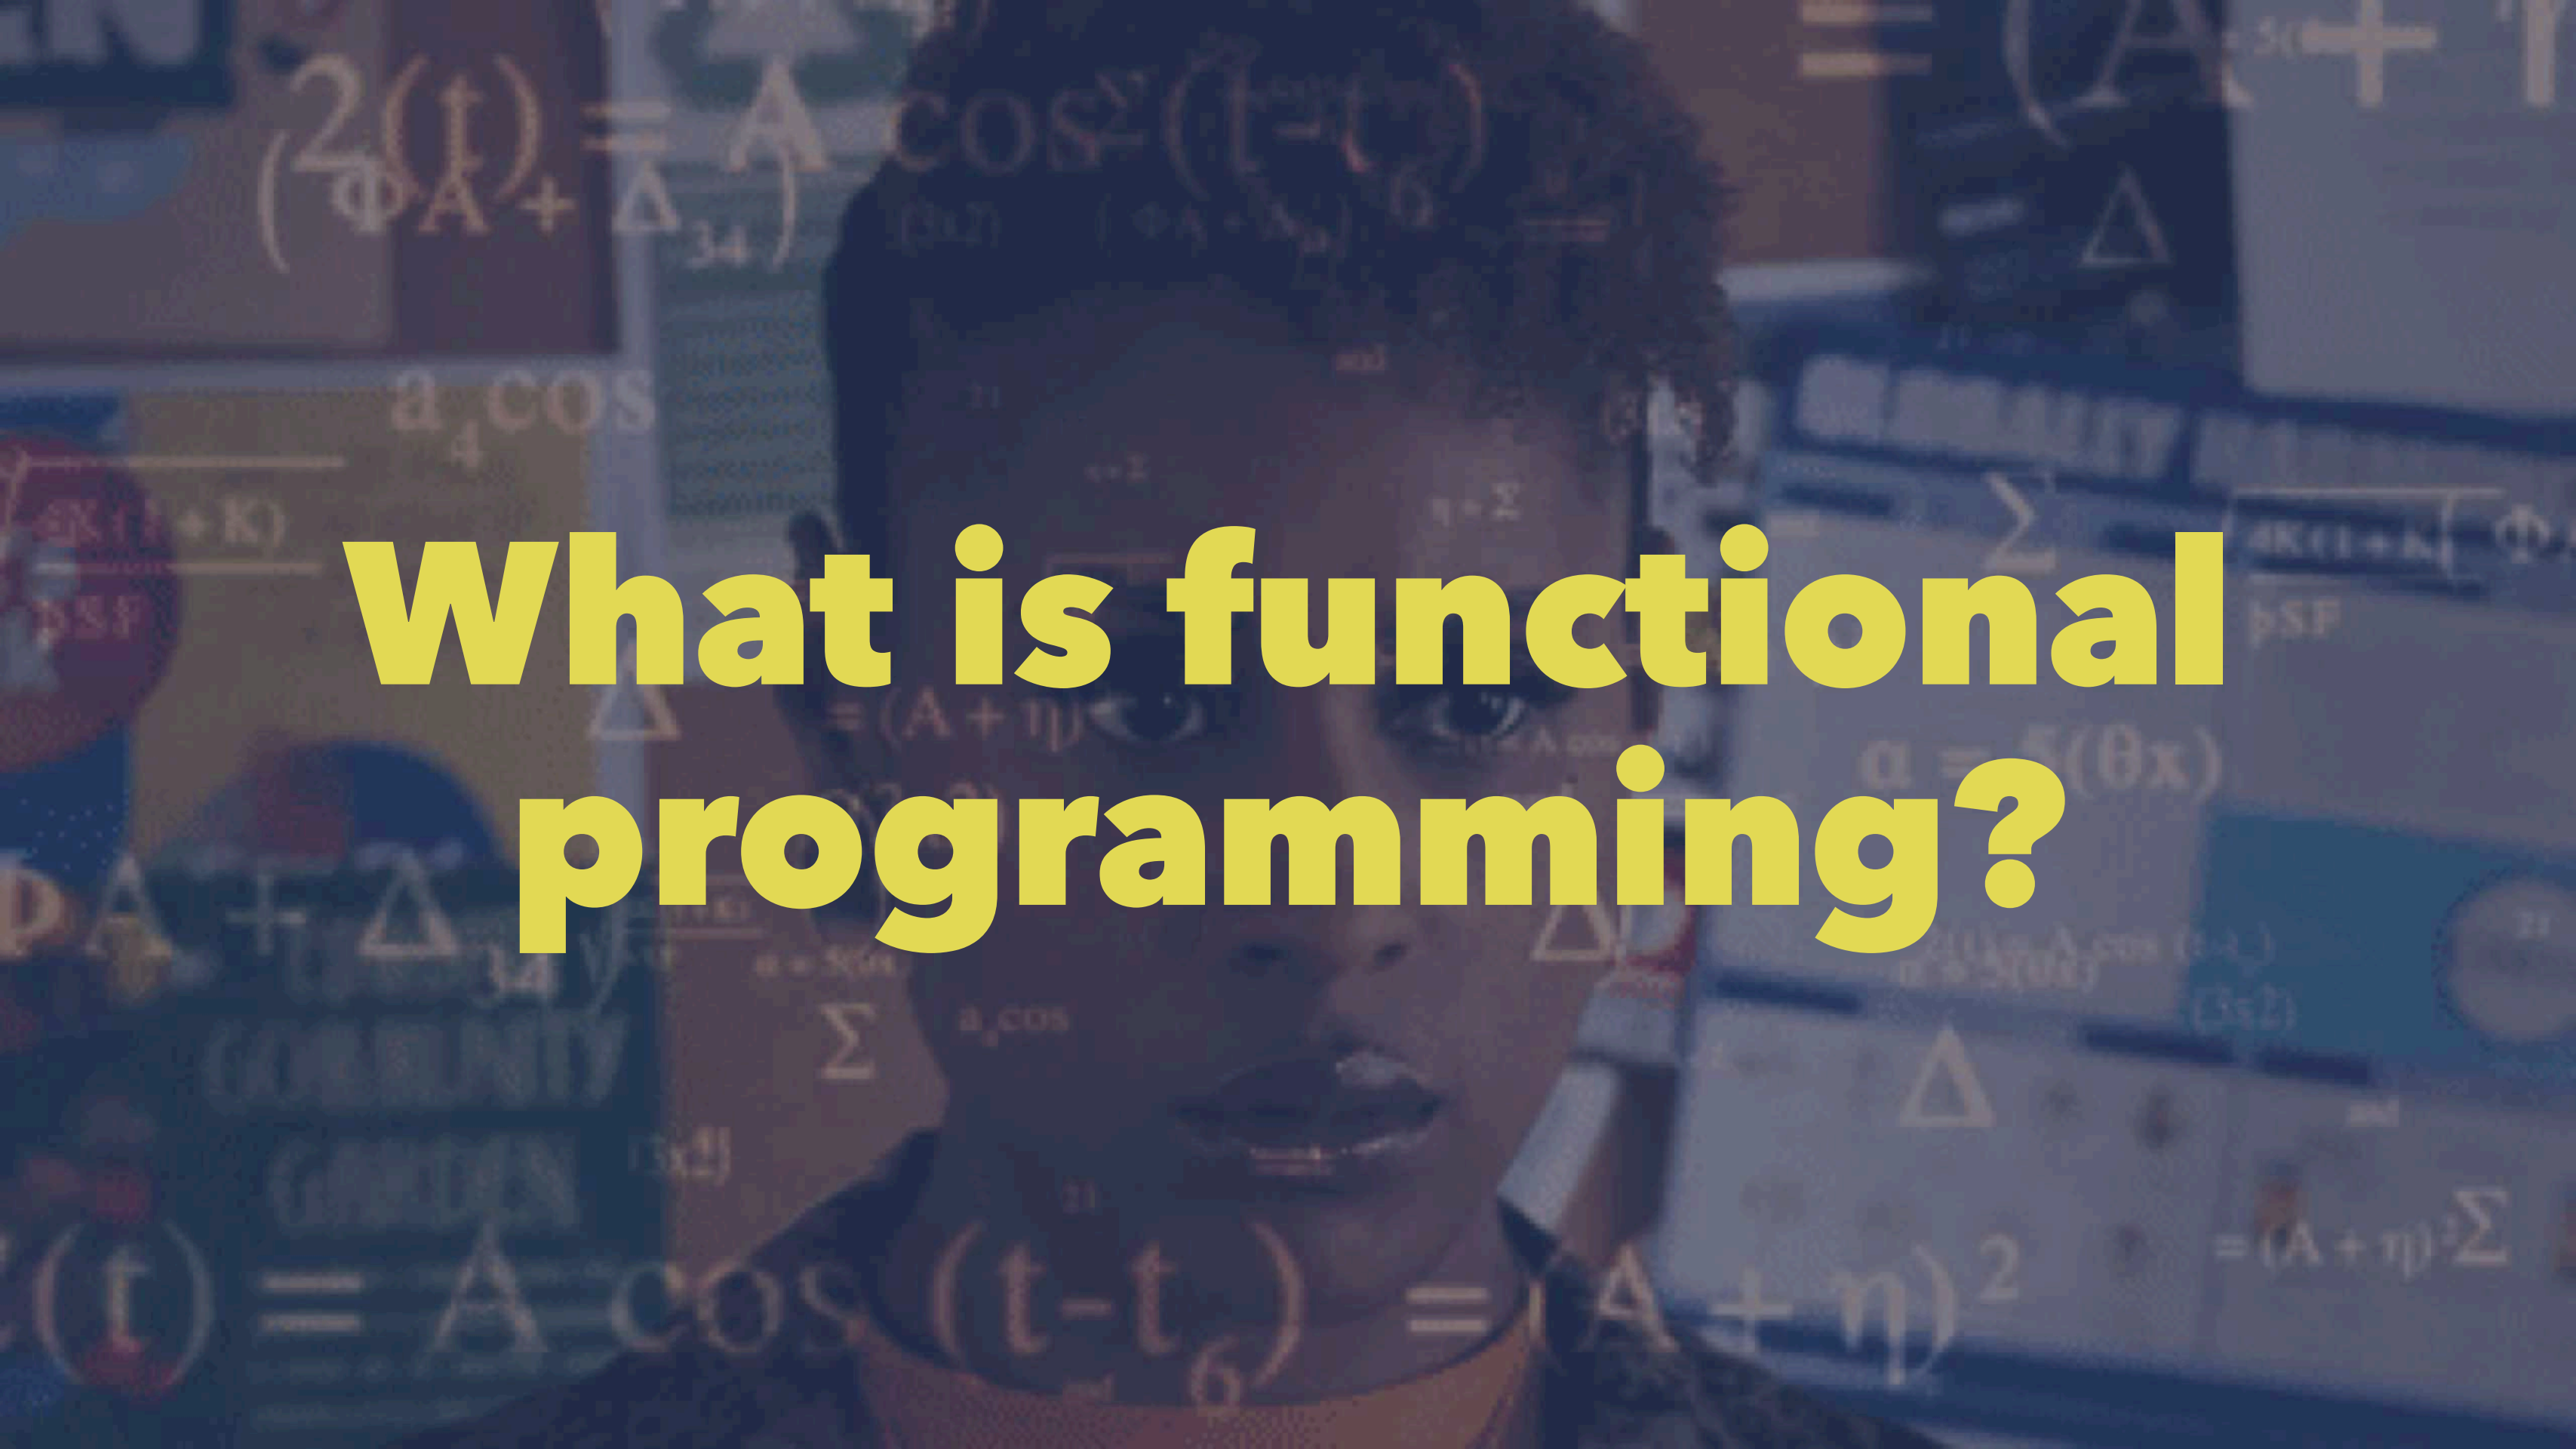
```
println(playerPosition.x)
```

- Case classes can be copied

```
val nextPosition =  
    playerPosition.copy(y = playerPosition.y + 1)
```

- They can be structurally compared

```
playerPosition == Position(18, 4) // => true
```




What is functional programming?

Functional programming is a *programming paradigm* that treats computation as the *evaluation of pure functions*.

(Or, as we call them, functions)

What are functions?

- In procedural programming, the term *function* is often wrongfully used to mean *procedure*, a way of binding a set of instructions to name and some arguments
- In FP, we treat functions as mere arrows between values of two sets
- Functions are first-class citizens: they can be passed around as arguments, returned from other functions and composed



A bit of history : Lambda-calculus

Lambda-calculus is the foundation of all functional languages. It can be described as the *most simple, universal programming language*

Lambda-calculus does not care about the computer. It abstracts away the computation steps under a purely expressive form, much like FP languages

Everything that can be computed by a Turing machine can be expressed in Lambda-calculus form¹.

¹ Look at the Church-Turing thesis

The syntax of λ -calculus

$expression :=$	$ variable$	(identifier)
	$ expression\ expression$	(application)
	$ \lambda variable. expression$	(abstraction)
	$ (expression)$	(grouping)

λ -calculus is tiny symbol manipulation framework.
It is entirely *expression-oriented*.

How does it work?

Consider this function:

$$f(x) = x^2$$

We can implement it in Scala like so:

```
def f(x: Int) = x * x
```

And expressed in λ -calculus a:

$$\lambda x. x^2$$

Abstractions in λ -calculus don't need to be named.

Lambda calculus introduced anonymous functions as values for the first time. I

In some programming languages, anonymous functions are still called lambda expressions today.

Abstractions in λ -calculus are mappings from one expression to another. They are *unary* (only one argument)

How would one write functions of two arguments in λ -calculus?

$$f(x, y) = xy$$

would become

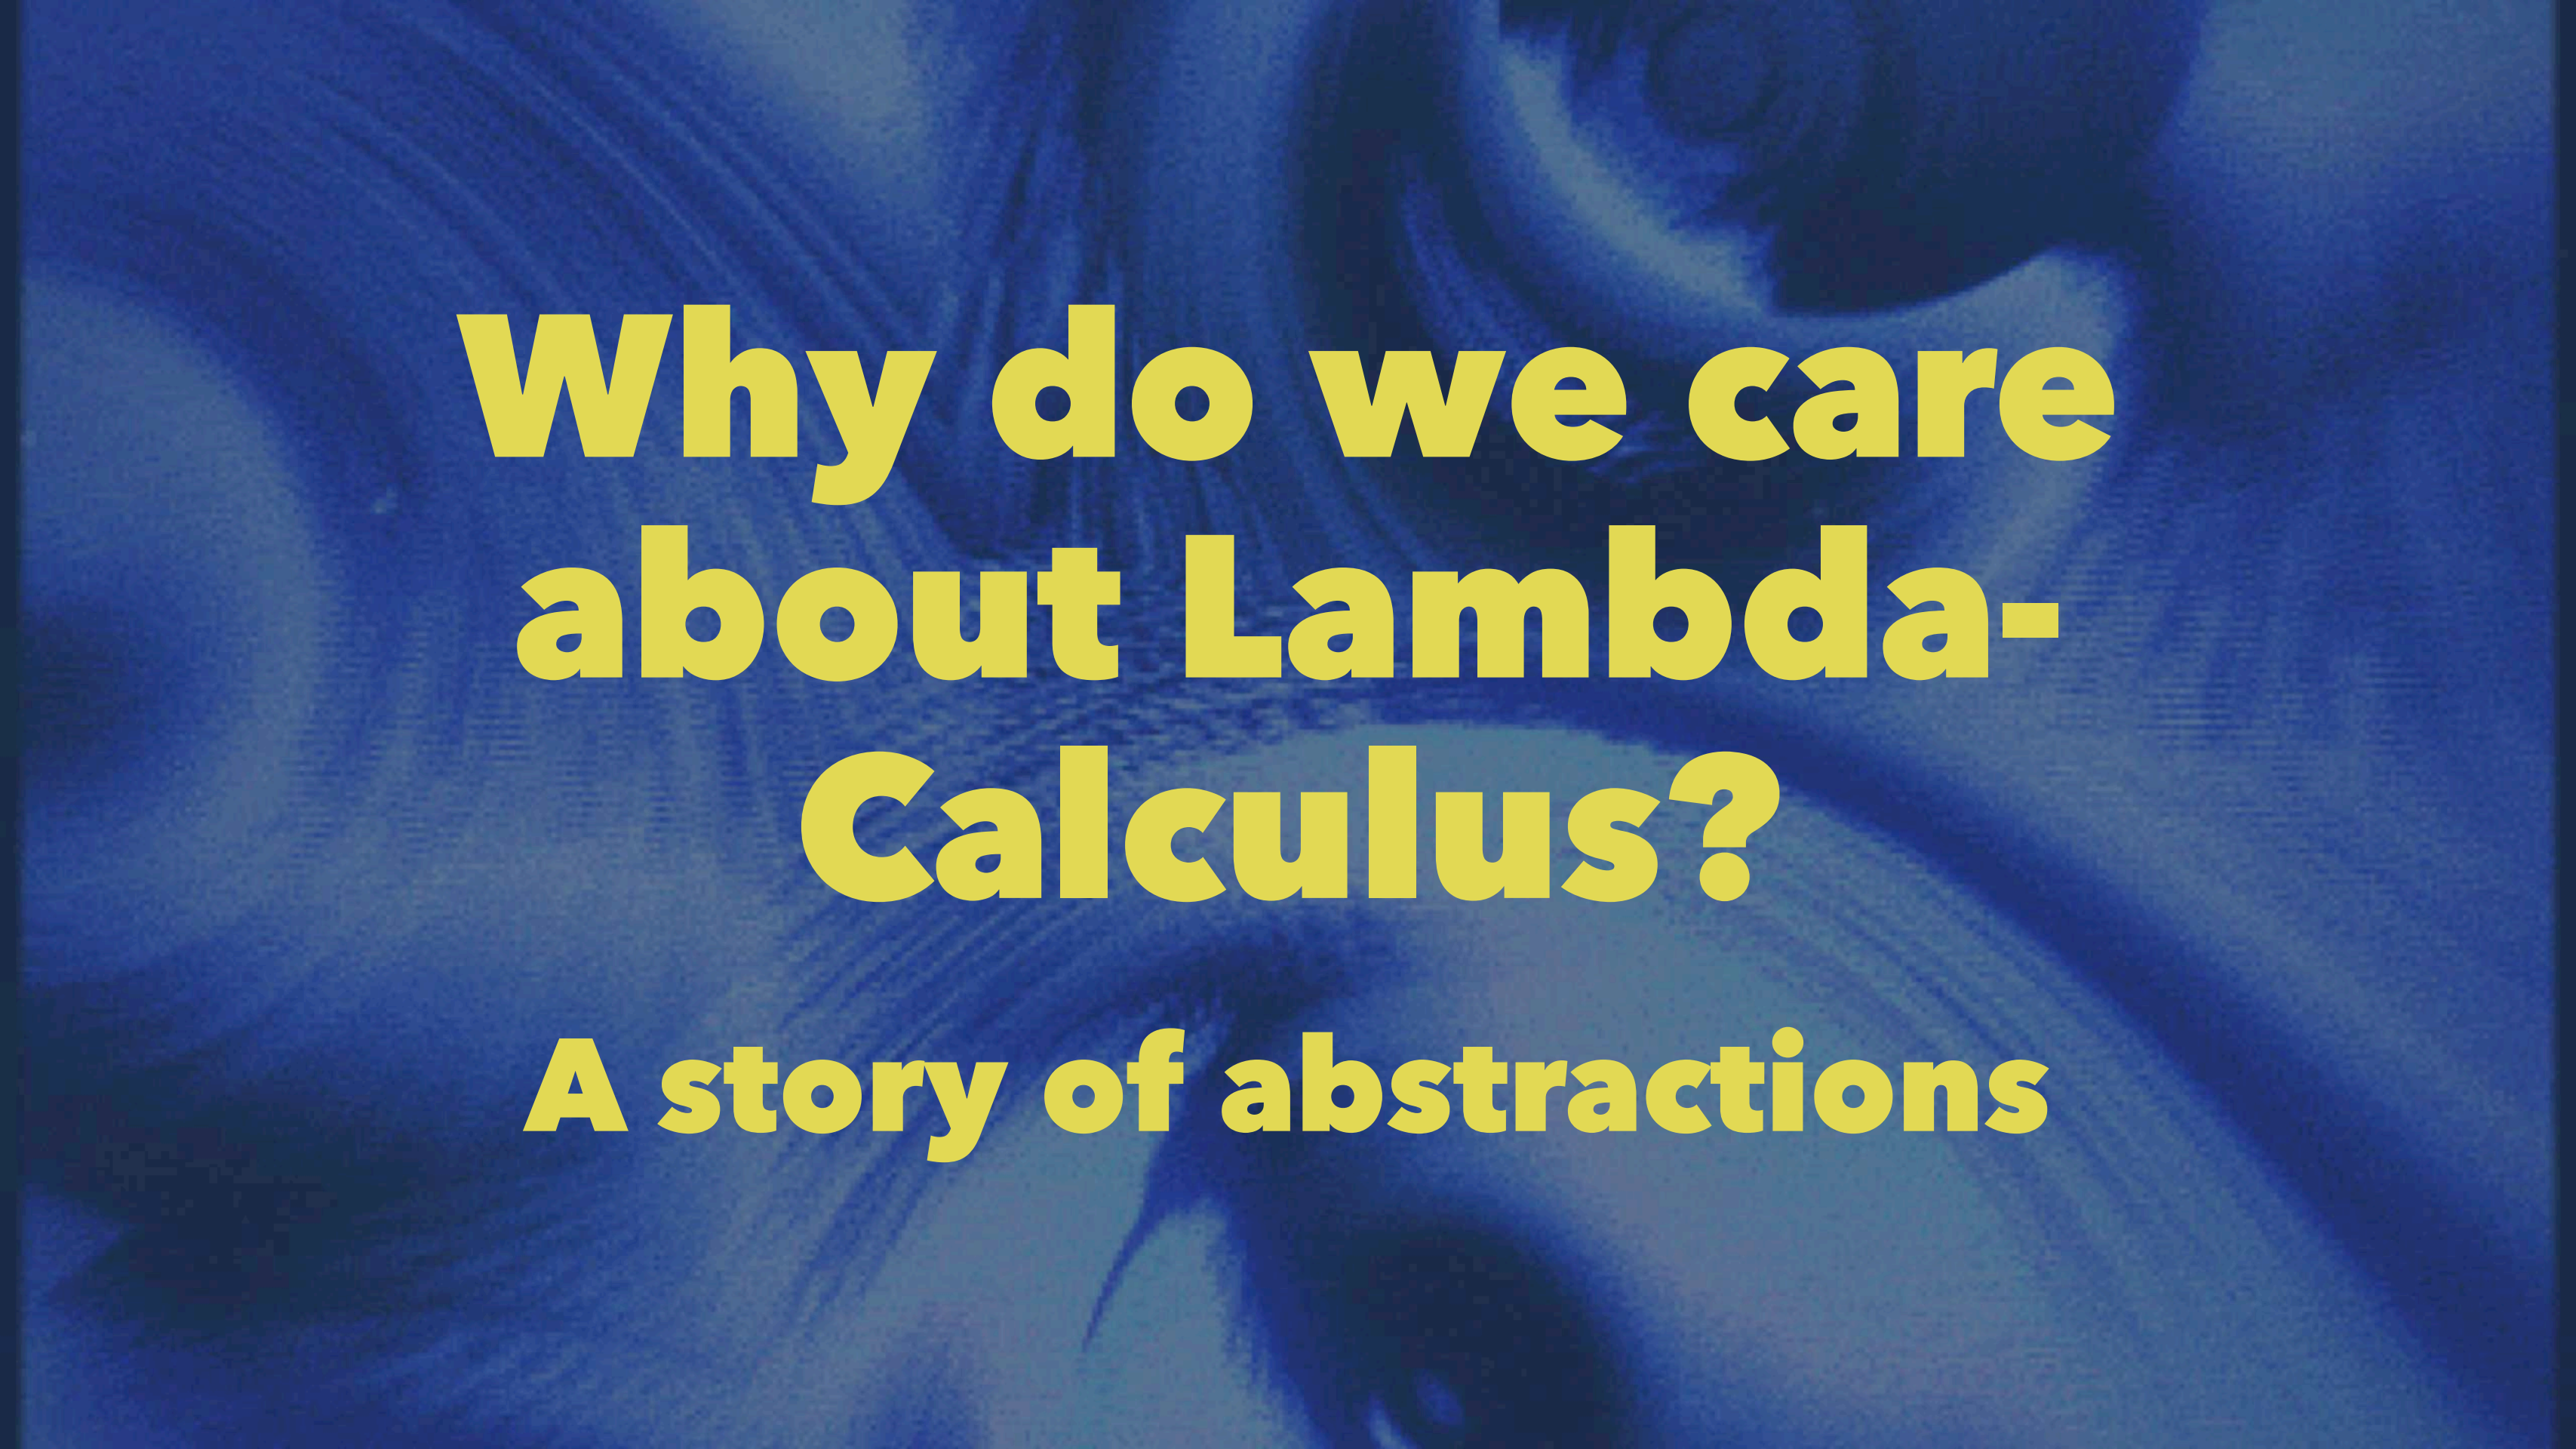
$$\lambda x. \lambda y. xy$$

A function of two arguments x and y can be expressed as a function of a single argument x , itself returning a function of a single argument y . This is called *currying*.

Most importantly, λ -calculus introduced *higher-order functions* : functions that can take return other functions or take them as argument. `map`, `filter`, `reduce` ... you already know some of them!

Some definitions

- α -conversion : $\lambda x.x = \lambda y.y$ (any variable can be renamed given it isn't free)
 - β -reduction : lambda-terms may be reduced to a simpler form through a succession of conversions, until we reach a form we cannot further simplify (β -normal form)
 - This is the foundation of program *evaluation*
- ⚠ Some terms cannot be reduced (halting problem)



Why do we care about Lambda- Calculus?

A story of abstractions

Turing machines are hypothetical devices that perform state-based computation on an infinite strip. Turing machines work *imperatively*.

On the other hand, Lambda-calculus is a purely abstract symbol rewriting framework. It's *expressive*.

Modern computers are built on the foundation of Turing machines.

One can usually categorize a programming language based on the level of abstraction it provides over the computer

- Modern computers (Turing-machine-inspired)
- Machine code
- Assembly languages
- Higher-level machine-centric languages (C),
- Higher-level imperative and object-oriented languages (C#, Java ...)
- Functional languages (LISP, Ocaml ...)
- Purely functional languages (Haskell, Idris ...)

Thinking beyond the machine

- Computers work imperatively. They execute instructions that change their state over time
- Functional programming languages operate on a higher level of abstraction : They don't care about the computer, they care very little for state, let alone time ... and it's great!
- Modern computers are more than capable of supporting these abstractions. Functional languages open up more possibilities for software engineering and developer productivity

Programming concepts that come from Lambda-calculus

- Lazy evaluation
- Lexical scoping and shadowing
- **Referential transparency**
- **Recursion**
- **Higher-order functions**
- Currying
- Closures
- ...

Defining pure functions

How to identify *pure functions* in an impure language such as Scala?

Referential transparency

- Referential transparency is one of the two key properties of pure functions
- A function is *referentially transparent* if you can substitute its *application* for its return value without changing the program's behavior

Example

```
def square(a: Int) = a * a  
val result = square(6)
```

And

```
val result = 36
```

are equivalent programs, which means that square is referentially transparent

We can also say it has no *side effects*.

Example

```
def square(a: Int) = {  
    println("Hi!")  
    a * a  
}  
val result = square(6)
```

is not equivalent to

```
val result = 36
```

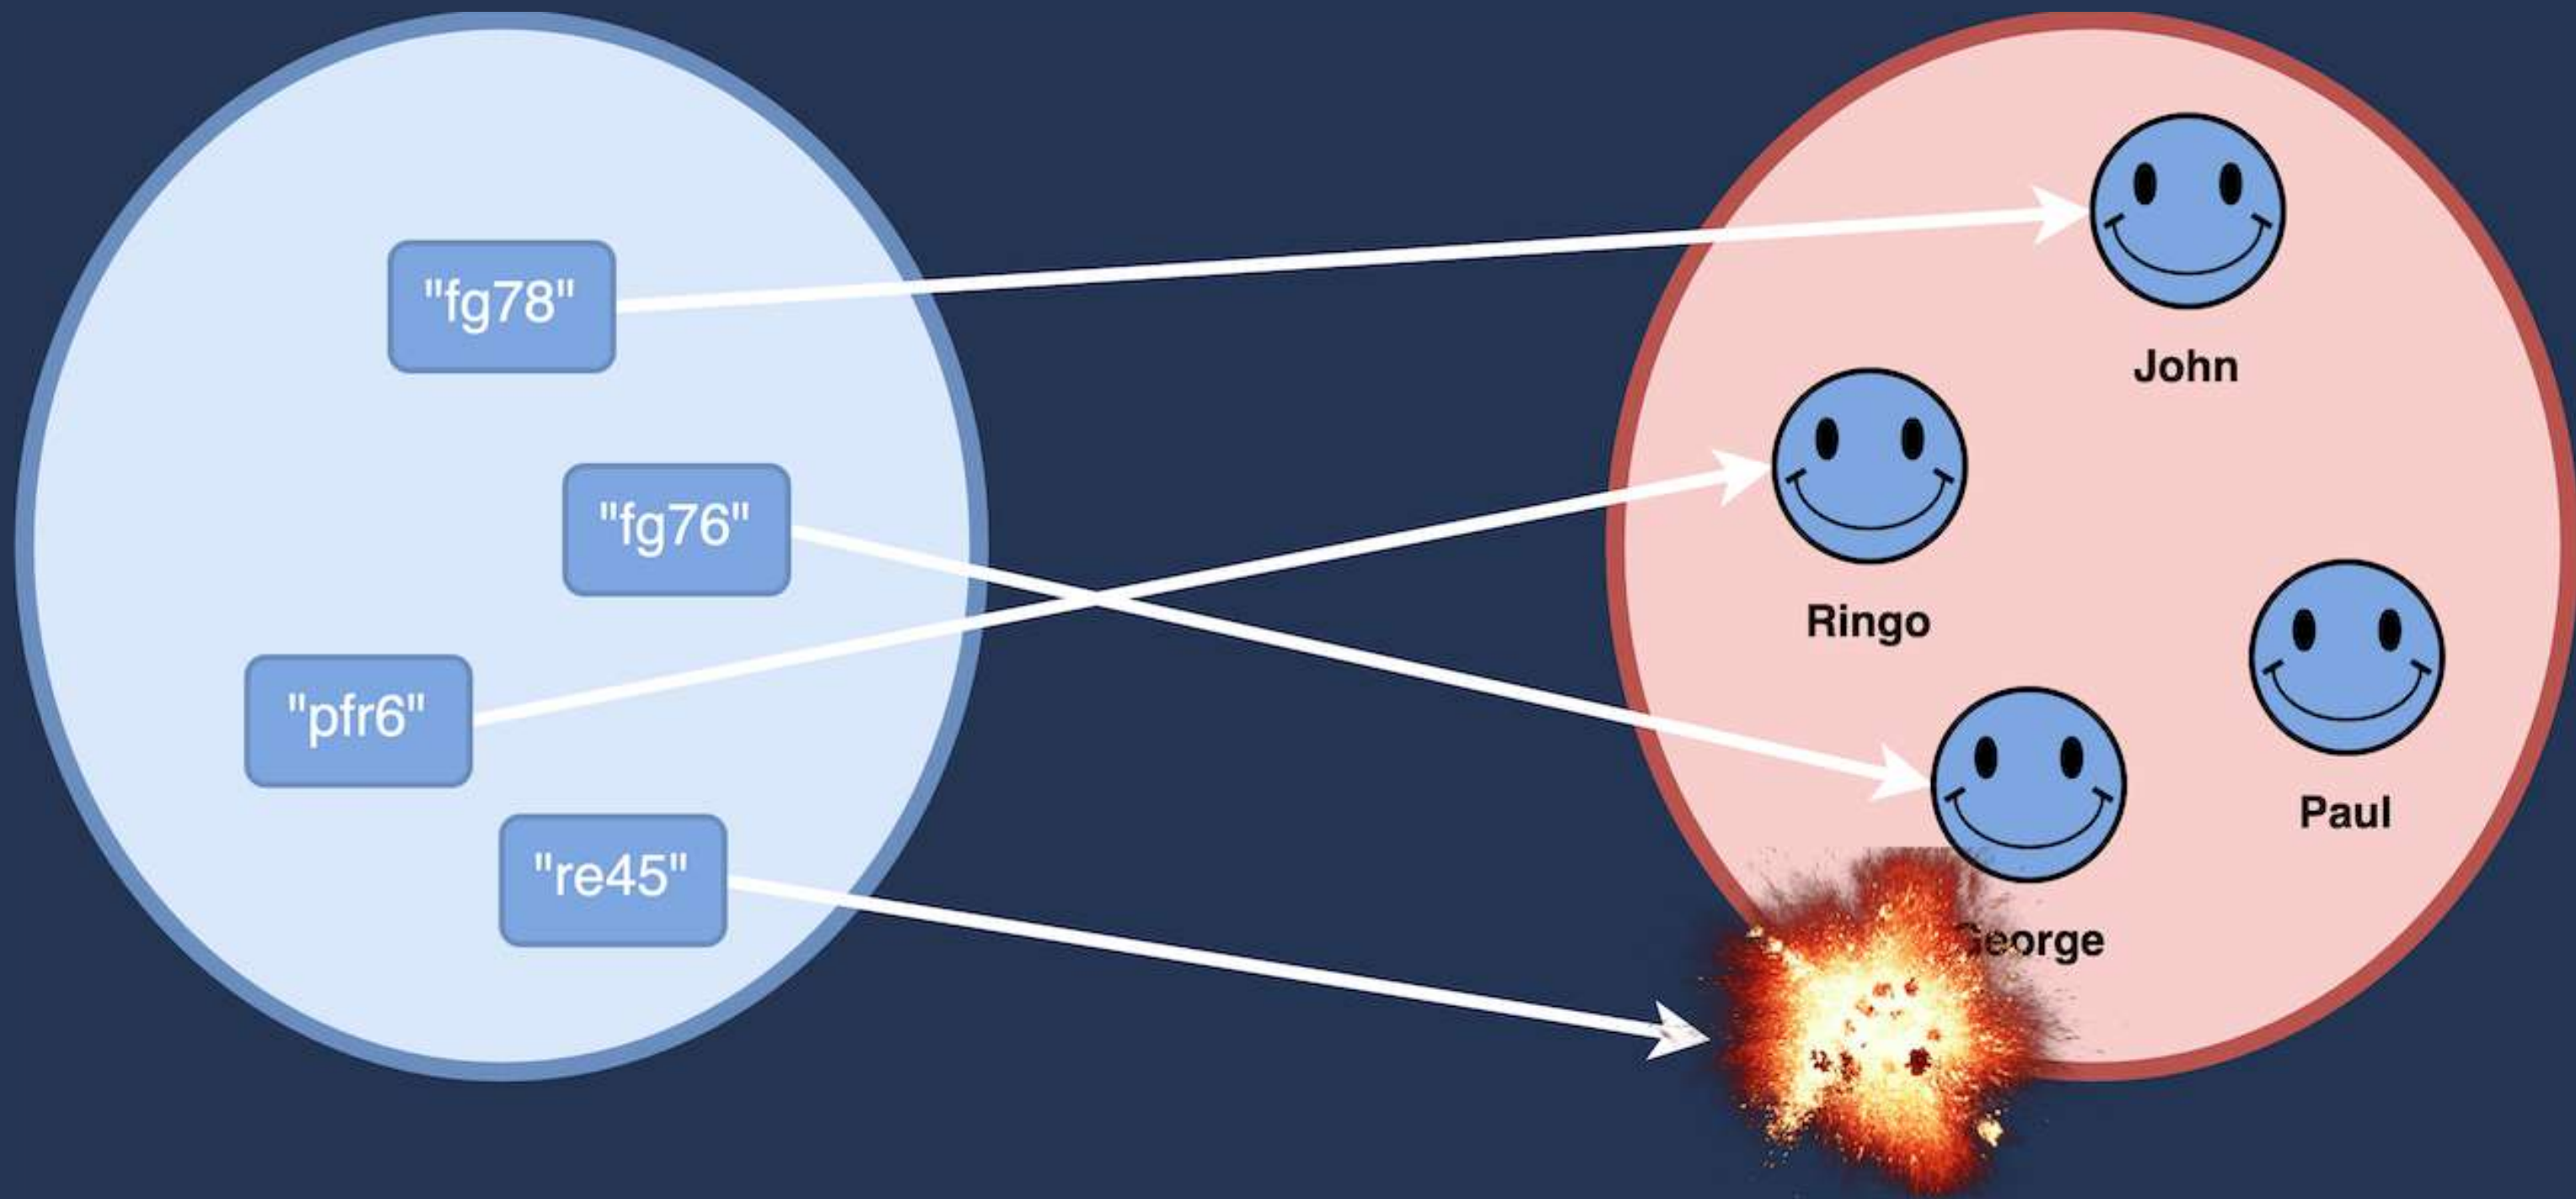
While the value of `result` is the same in both cases, the runtime behaviors of the two programs is not the same.

Why should I care? Benefits of referential transparency

- Referentially transparent functions are trivial to test
- They are easy to execute in parallel
- Predictability! You know exactly how your program will run
- Memoization : if a function is guaranteed to return always the same result for some input, it can be memoized (cached) without risks

Totality

- Totality is the other key characteristic of pure functions
- A function $A \Rightarrow B$ must absolutely yield a value of B for *any* value A
- Partial functions (functions that operate on a subset of A) are unsafe
- As a rule of thumb, *if it compiles, it must always work*
 - Exceptions break totality and should be avoided in most cases



A short definition for pure functions

**A function is pure if, and only if, it
is both total and referentially
transparent**

Spotting impure functions in the wild

Is this function pure?

```
def max(a: Int, b: Int) =  
    if (b >= a) b  
    else a
```

Remember to ask yourself:

- Does it have side effects?
- Does it return a value for every possible argument?
- Is the return value defined by the arguments only?

What about this one?

```
var counter = 0
def getName: String = {
  counter += 1
  "Jessica"
}
```

And this one?

```
def mean(numbers: List[Int]): Float = {  
    println("Computing mean")  
    numbers.sum.toFloat / numbers.length  
}
```

This one is trickier

```
def mean(numbers: List[Int]): Float = {  
    numbers.sum.toFloat / numbers.length  
}
```

- Sometimes, identifying an impure function isn't as straightforward.
- How could we make it safer?

What bout his last one?

```
def getSecret(user: User): Secret =  
    if (user.isAdmin()) {  
        secret  
    }  
    else {  
        throw new Exception("You don't have access")  
    }
```

From now on, we will only use pure functions

Wait, how am I supposed to do stuff if everything is forbidden?

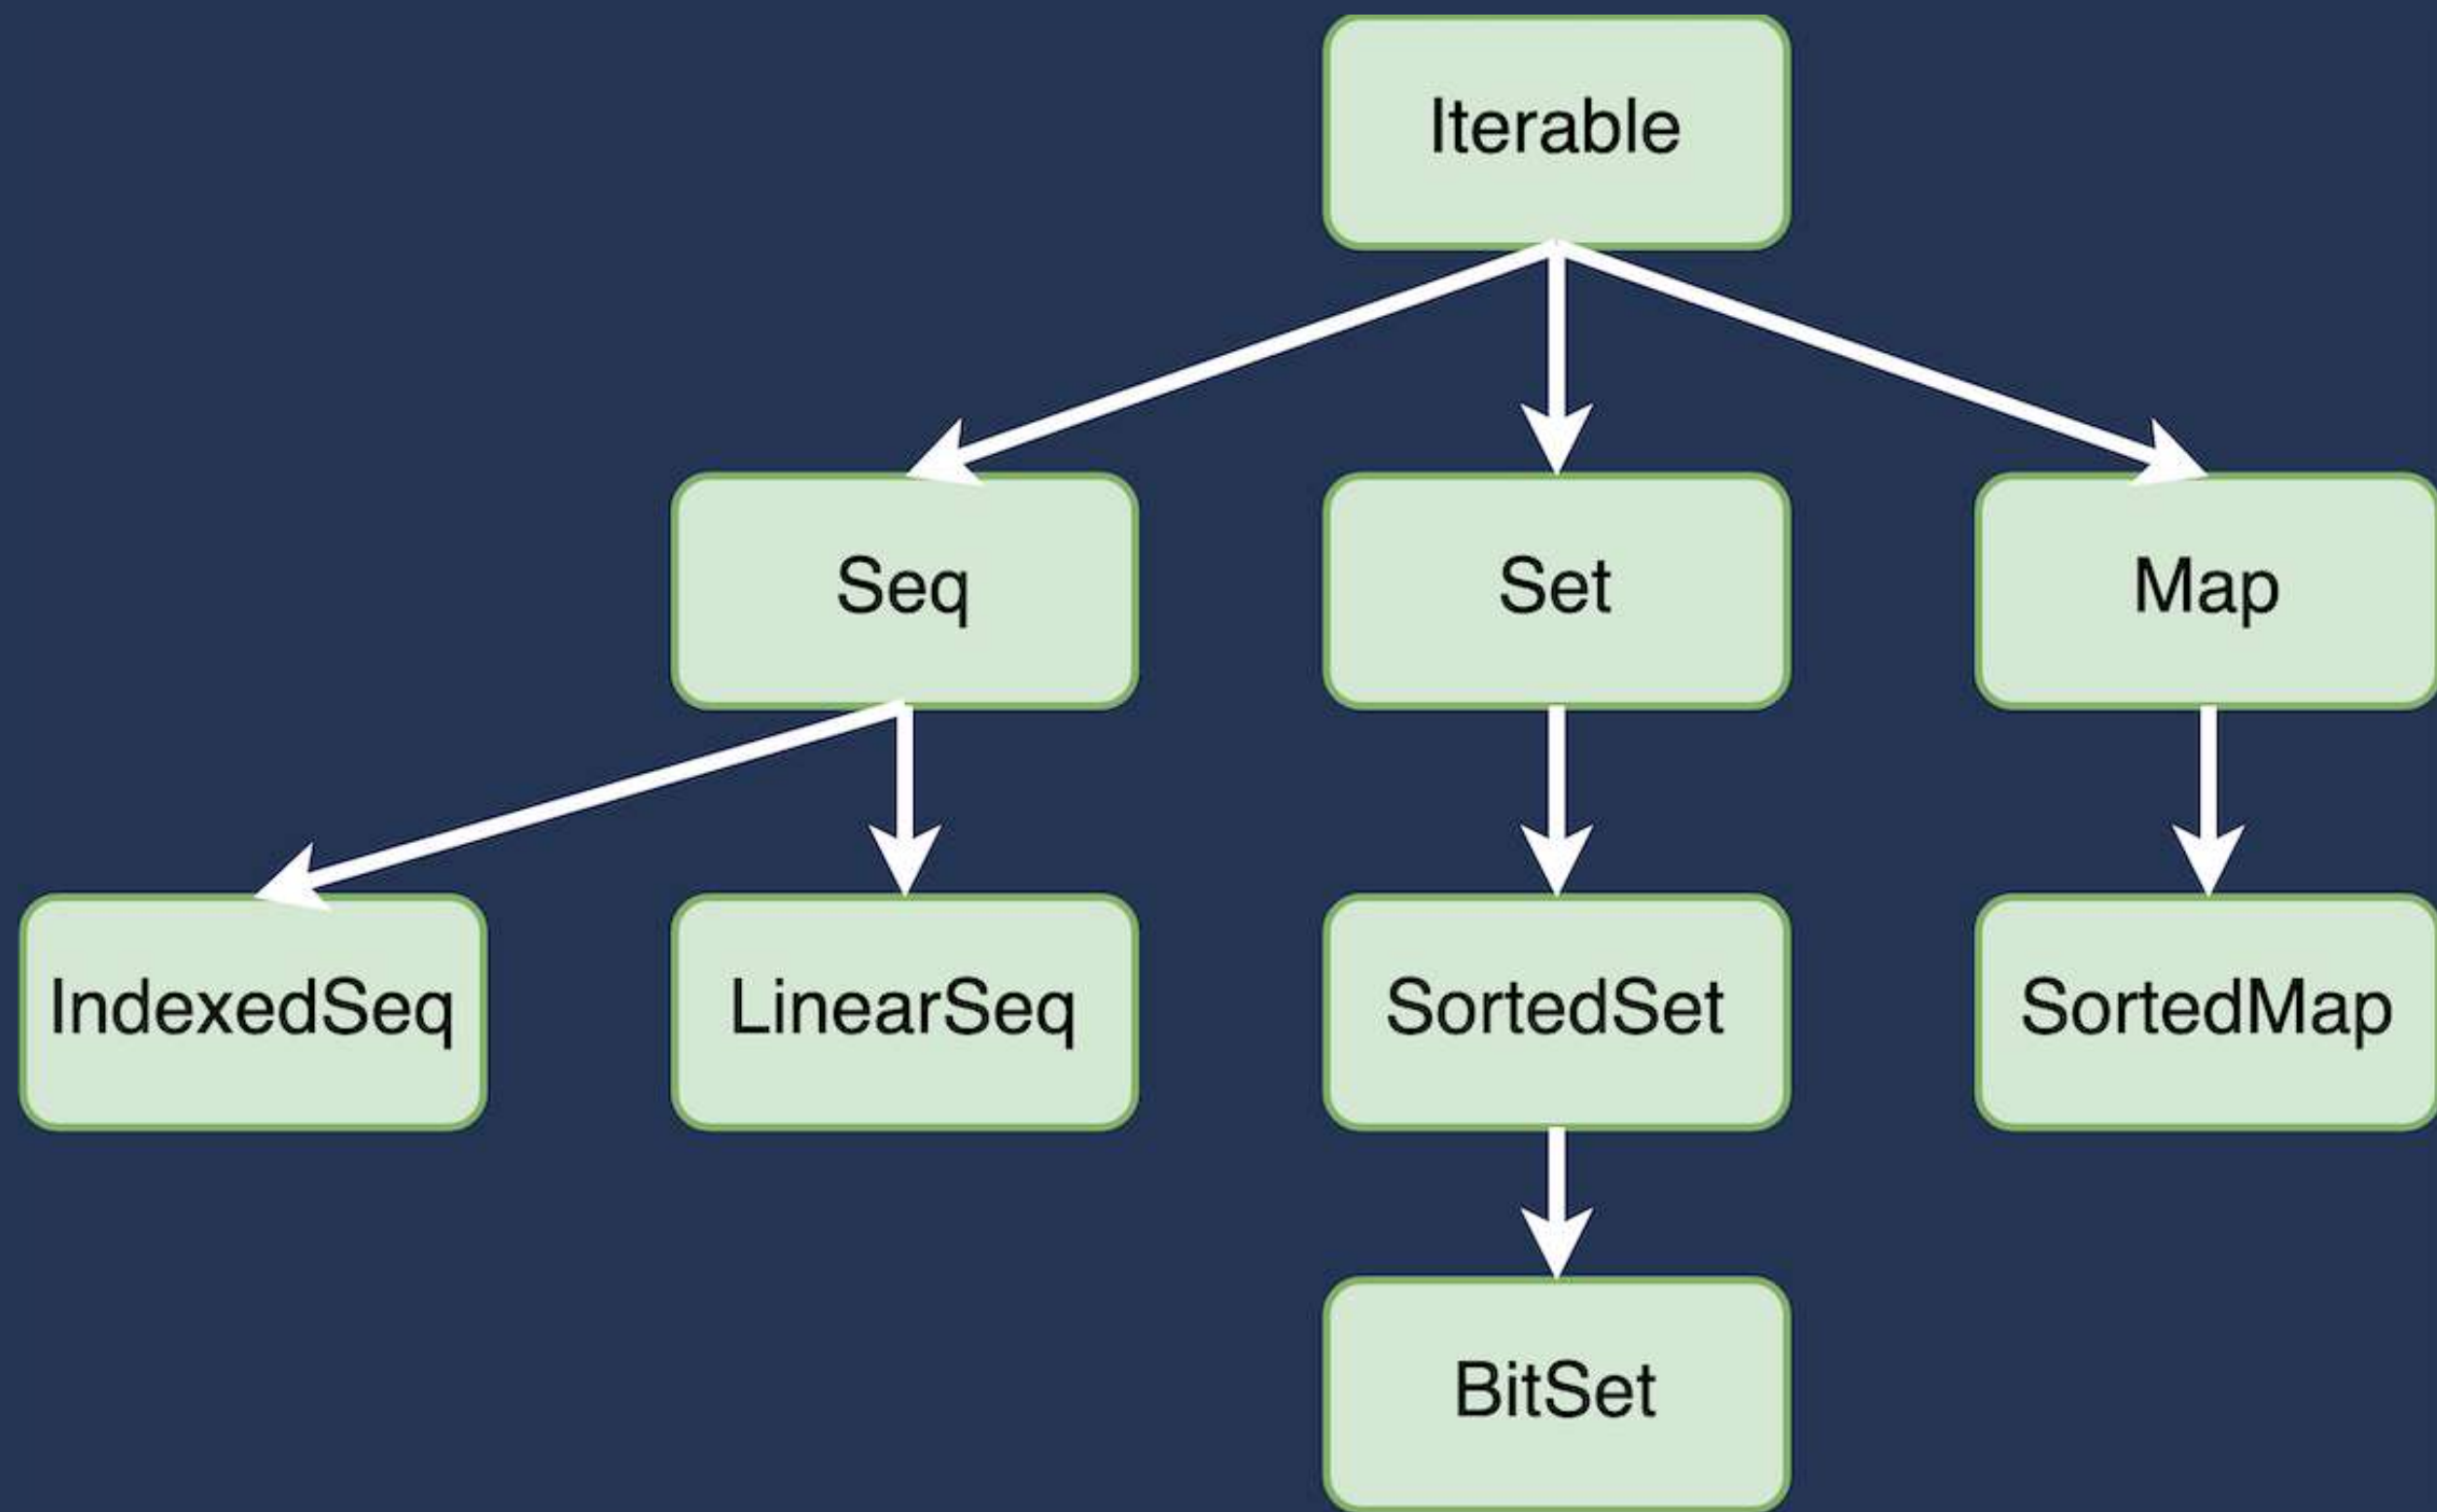
Functional programmers have strategies

- Persistent Data Structures
- Algebraic Data Types
- Pattern matching
- IO monads

Scala collections

Collections.scala

- Scala provides several collections that can model lists of elements, ranges, sets of unique elements, maps ...
- Scala's collections are persistent by default, which means you can't modify them in place
- Scala's collections are organized into a hierarchy. Different concrete data structures can be manipulated through common interfaces



Which do I need?


- Is your code's performance critical?
 - Yes: check performance characteristics of all data structures in the documentation
 - No:
 - If you need a list of possibly duplicated elements, use `List`
 - If you need a list of unique elements, use `Set`
 - If you need to identify elements using a key, use a `Map`
 - If you're writing a library, it's a good idea to abstract over the collection type using an interface

Working with Lists

- `map`: If you have `List[A]` and `A => B`, you get `List[B]`
- `flatMap`: If you have `List[A]` and `A => List[B]`, you get `List[B]`
- `filter`: If you have `List[A]` and `A => Boolean`, you get a subset of your original list
- `flatten`: If you have `List[List[A]]`, then you get `List[A]`
- `headOption` : A safe way of retrieving the first element
- `lift` : A safe way of retrieving any element
- `sum` (on a numerical list only)

Structural sharing


```
val a = List[User](john, ringo, george)
```



```
val paul = User("Paul", "McCartney")
```



```
val fabFour = paul :: a
```



A couple things to remember

- Lists are good at prepending (constant time) but not so much at appending (linear time)
- Persistent data structures are space efficient : since every value is reputed immutable, they can be shared by reference across multiple structures

A red brick house with a gabled roof and several windows, rendered in a pixelated, blocky style. It sits on a grey surface against a grey background.

Building new types

Domain modeling using Algebraic Data Types

ADT.scala

Product Types


Tuples and case classes can encapsulate the values of several types into one single value :

```
val cathy: (String, Int, Role) =  
    ("Cathy", 38, Role.Manager)
```

The cartesian product

Product types are called that way because they can be seen as the cartesian product of several sets :

- Types, such as `Int` and `Boolean` are sets
- A tuple `(Int, Boolean)` can match any `Int` with any `Boolean`
- The total number of distinct values this tuple can have will be
 - The number of values `Int` can take
 - Multiplied by the number of values `Boolean` can take (2)



Case classes are product types

Case classes can be seen as tuples with labels

```
case class User(  
  firstName: String,  
  lastName: String  
)
```

can be expressed as

```
type User = (String, String)
```

However, tuples are ambiguous. Which is the `firstName` in that case ?

Keep tuples for very simple values and internal use only.

Sum types

- Sum types can hold values coming from exactly one of several fixed types
- They are sometimes called *variants*, *disjoint unions* or *discriminated unions*
- They can be recursive

Common subtypes you might already know

- `Option[A]`: One of `Some[A]` or `None`
- `Either[L, R]`: One of `Left[L]` or `Right[R]`
- `List[A]`: One of `Nil` or `::(A, List[A])`

Lists are recursive sum types!

The `List` structure, a linked list, is sum type of two *terms* :

```
sealed trait List[+A]  
// A Nil term representing an empty list  
case object Nil extends List[Nothing]  
  
// A :: term representing an element  
//and a pointer to the rest of the list  
case class ::[A](head: A, tail: List[A]) extends List[A]
```

This is how we would build lists using that encoding:

```
Nil                                // An empty List
::(1, Nil)                        // A List of one element
::(1, ::(2, Nil))                 // A list of two elements
::(1, ::(2, ::(3, Nil)))          // And so on ...
```

Luckily, Scala provides some syntactic sugar for us:

```
val nb = 1 :: 2 :: 3 :: 4 :: Nil
```

Algebraic Data Types are the main tool for data modeling in Scala

A sum type for our Tic Tac Toe

We can model a *move* in the Tic Tac Toe game using a sum type

```
sealed trait Move  
case object X extends Move  
case object O extends Move
```

NB: Case objects are singleton objects. You can see them as sets of exactly one element



Pattern matching

PatternMatching.scala

Pattern matching

- Pattern matching is a pure mapping between a *pattern of expressions* on the left side, and an expression on the right side.
- It can deconstruct case classes and match their individual fields
- It can have arbitrary guards (if statements)
- Exhaustivity can be enforced at compile time

How does that help us?

- Persistent data structures give us
 - Complex data transformations without the need for shared mutable state and its risk
- ADTs coupled with pattern matching give us
 - Powerful data modeling
 - Total functions
 - Safety from the most common runtime exceptions

General strategies for building total functions

In order to make a function total, you generally have two possible approaches:

- Restrict the input of the function (or *domain*) so that your function can't be called with input you can't deal with
- Expand the output of the function (or *codomain*) so that you can produce a result in every case

Expanding the output

Expanding the output generally means turning your return type from `A` to `Option[A]`, `Either[Error, A]` or another structure.

Instead of returning values from the smaller set `A`, you operate on a bigger set, the set of `A + all the possible edge-cases`.

You essentially force your caller to deal with the edge cases somewhere outside.

Restricting the input

You can also change the types of your arguments to accept a smaller set of values. This way, you ensure your function cannot be called with values you can't handle.

For instance, instead of accepting any `String`, you can use a sum type as your input type.

Summing up

Scala let's you perform the most common computing tasks without relying on side-effects.

Writing total functions means you don't need Exceptions to deal with edge-cases, hence no crashes at runtime.

Writing referentially transparent functions means you can refactor your code safely. It also means *no lies* : your types should always tell the truth



Friends don't lie.

**But still,
what if I *need* that side-effect?**

Introducing IO monads

SideEffects.scala

IO monads

- The IO monad comes from Haskell, a purely functional language where side effects are strictly forbidden
- An `IO[A]` represents a *lazy* computation that has some side effects and eventually yields a value of type `A`
- IOs can be transformed with `map` and composed with `flatMap` just like `Option` and `List`

Why do we care?

- The #1 reason: **functions signatures are contract*. When you perform a side-effect without signaling it, you break the contract. (aka. *Friends don't lie*)
- IOs reduce the need for further documentation and the risk of forgetting to catch an exception
- They are referentially transparent, so you can define them in any order
- In a FP team, other functions should be assumed to be pure



Daniel Brice

@fried_brice



Show me your implementations and conceal your types, and I shall continue to be mystified. Show me your types, and I won't usually need your implementations; they'll be obvious.

[Traduire le Tweet](#)

12:29 AM · 3 sept. 2019 · [Twitter Web App](#)

4 Retweets **31** J'aime

Other benefits from IO

- They can raise and recover exceptions
- They can easily be executed in parallel
- They form a large ecosystem that makes concurrent applications easier
 - Concurrency constructs from Cats Effect such as `Ref`, `MVar`...
 - Reactive streams with `fs2`
 - Purely functional database access with `Doobie`
 - Highly performant HTTP APIs with `Http4s`
 - ...

IO is a Monad. What is that?

A Monad $M[A]$ is an immutable data structure used to describe the computation of one or more values of type A

- Values of type A can be turned into *monadic values* $M[A]$
- Monads can chain subsequent computations (i.e solve big problems out of smaller problems)
- They describe some *functional effect*: `Option` describes optionality, `Either` describes failure ...

Monads $M[A]$ have essentially four elements:

- A `pure` method that wraps A into the type constructor $M[A]$
- A `map` method of the form $M[A] \Rightarrow (A \Rightarrow B) \Rightarrow M[B]$,
which is the definition of a *functor*
- A `flatMap` method of the form
 $M[A] \Rightarrow (A \Rightarrow M[B]) \Rightarrow M[B]$
- Laws called *identities*

**You already know how to use
them**

	pure	the functor	flatMap
M[A]	A => M[A]	M[A] => (A => B) => M[B]	M[A] => (A => M[B]) => M[B]
Option[String]	Some("Hello")	Some("Hello") .map(_.toUpperCase) // Some("HELLO")	Some("Hello") .flatMap(_ => Some("World")) // Some("World")
List[String]	List("foo", "bar")	List("Foo", "Bar") .map(_.length) // List(3, 3)	List("ab", "cd") .flatMap(_.toList) // List('a', 'b', 'c', 'd')
Either[A, B] IO[A] Reader State ...		They all have the same shape	

Why monads?

- Monads are used to encode *functional effects* such as optionality or failure, in a way that **makes the effect appear in the signature**
(Remember, signatures must always tell the truth)
- The essence of a monad is *composition*, the ability to chain dependent computations
 - All monads can be used in a *for-comprehension*

Because they can encode all sorts of effects using the same signatures, monads can encode features that would otherwise require specific syntax at the language level.



Diogo (f a) Castro

@dfacastro



Welcome to Haskell, we got list comprehensions ($>>=$), optional chaining ($>>=$) , async/await ($>>=$), if err != nil ($>>=$), function composition($>>=$), mutable state ($>>=$),

Your first, purely functional, Scala application

Building a functional Tic Tac Toe

Let's use what we've learn to build a command-line Tic Tac Toe that can be played on the network!

- Tic Tac Toe can be modeled as a *finite state machine*
- Possible states = a State ADT
- State transitions = pure function of State \Rightarrow State
- The board will be a matrix of *moves*, modeled as a two-dimensional List
- The user interface and server can be separated from the core logic using higher-order functions and IO monads